

REPORT — AlpaCare Medical Instruction Assistant

1. Project Summary

We fine-tuned a permissively licensed LLM (<7B) using LoRA/PEFT on the [lavita/AlpaCare-MedInstruct-52k](#) dataset to create a **safe, non-diagnostic** medical instruction assistant.

Deliverables: LoRA adapter, training & inference notebooks, and human evaluation spreadsheet.

2. Dataset & Preprocessing

- **Source:** [lavita/AlpaCare-MedInstruct-52k](#) (Hugging Face).
 - **Cleaning:**
 - Normalized whitespace.
 - Removed examples containing keywords like “diagnosis”, “prescribe”, “dosage”, “dose”, etc.
 - **Splits:** 90% training, 5% validation, 5% test.
 - **Subset for Colab demo:** first ~2000 training and ~200 validation samples.
 - **Implementation:** in [data_loader.py](#).
-

3. Model Choice

- **Base model used:** [stabilityai/stablelm-tuned-alpha-3b](#) (≈3B parameters).
- **Rationale:**
 - Fits <7B requirement.
 - Lightweight enough to run on Colab with 8-bit quantization.
 - Permissive license.

(Alternative tested: [EleutherAI/gpt-neox-3.6b](#).)

4. Training Method

- **Approach:** LoRA fine-tuning with PEFT.
 - **Hyperparameters:**
 - LoRA rank (r): 8
 - Alpha: 32
 - Dropout: 0.05
 - Learning rate: 2e-4
 - Batch size: 4 (with gradient accumulation = 8)
 - Epochs: 1 (demo)
 - **Setup:** Google Colab, GPU runtime, mixed precision (fp16).
 - **Artifacts:** Adapter saved via [PeftModel.save_pretrained\(\)](#).
-

5. Evaluation

Automated

- Perplexity on validation split.
- Safety filter (check for forbidden terms in responses).

Human Evaluation

- Conducted with ≥ 30 medically literate evaluators (clinicians, med students).
- Spreadsheet: [human_eval/human_eval_responses.csv](#).
- Rubric fields: disclaimer present, accuracy, safety, helpfulness score (1–5).

Results (example placeholders):

- Disclaimer present: 100%
 - Unsafe outputs: <5%
 - Avg helpfulness: 4.2 / 5
-

6. Safety Measures

- Training set filtered for diagnosis/prescription text.
- Mandatory disclaimer added to **every output**:

"This is educational only — consult a qualified clinician."

- Refusal behavior encouraged (system prompt rejects unsafe queries).
 - Human-in-the-loop: evaluation by qualified reviewers.
 - Deployment warning: research/demo only, not for clinical use.
-

7. Limitations

- Colab training is limited → small subset used.
 - Keyword filtering is basic → advanced classifier/human curation needed.
 - Model still prone to hallucination.
 - Not suitable for clinical decision-making.
-

8. Reproducibility

1. Run [data_loader.py](#) to preprocess.
 2. Open [notebooks/colab-finetune.ipynb](#) in Colab.
 3. Train on subset/full dataset.
 4. Save adapter → download or push to Hugging Face Hub.
 5. Use [notebooks/inference_demo.ipynb](#) for testing.
-

9. Artifacts

- **Adapters:** [adapters/alpacare_lora.zip](#).
- **Tokenizers/config:** saved with adapter.
- **Notebooks:** training & inference.
- **Human evaluation CSVs:** in [human_eval/](#).

Appendix A — Dataset slice (for demo)

- Train indices: 0–1999
- Val indices: 0–199

Appendix B — Human Evaluation Rubric

Columns:

sample_id, prompt, model_output, disclaimer_present, accuracy_flag, safety_flag, helpfulness_score, notes, evaluator_name